

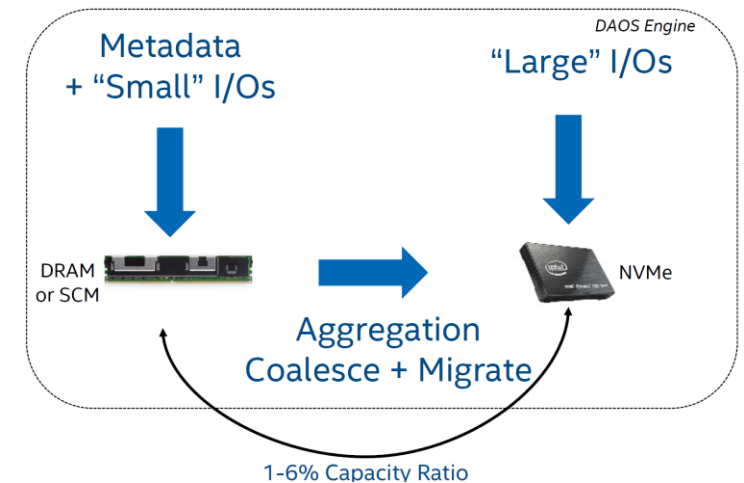
DAOS

HPC Object Store



VOS – Versioning Object Store

- DAOS is implemented as a VOS
 - Epochs and transactions built into the storage
- Updates create new version of the *value* in the kv
- Epochs used to provide consistency on updates
- Can use transactional APIs to group data updates atomically in applications
- Small I/O operations stay at the metadata level
 - Larger I/O operations hit the main storage
 - Dynamically managed



Non-NVRAM DAOS

- Path to solution with no persistent memory
 - Metadata in RAM
 - Write Ahead Log (WAL) on NVMe
 - Synchronous WAL updates
 - Asynchronous WAL checkpointing
- Reads faster as from DRAM
- Currently limited by DRAM capacity (async checkpointing done but not leveraged for lower resident set yet)
 - Restricts the size of DAOS or requires expensive DRAM setup
 - Allows a wide range of hardware for DAOS servers
- Google are using it for high performance storage in GCP
 - ParallelStore
 - Volatile deployment, Lustre alternative
 - “Parallelstore is based on [Intel DAOS](#) and delivers up to 6.3x greater read throughput performance compared to competitive Lustre scratch offerings.”

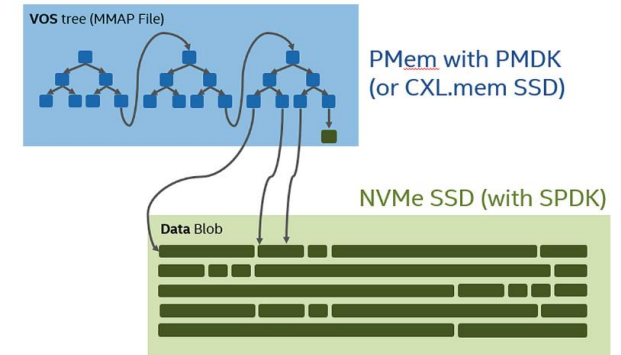


Fig. 2. DAOS Backend using Persistent Memory.




Fig. 3. DAOS Backend using Volatile Memory.

DAOS Beyond Persistent Memory: Architecture and Initial Performance Results

https://doi.org/10.1007/978-3-031-40843-4_26

Using DAOS

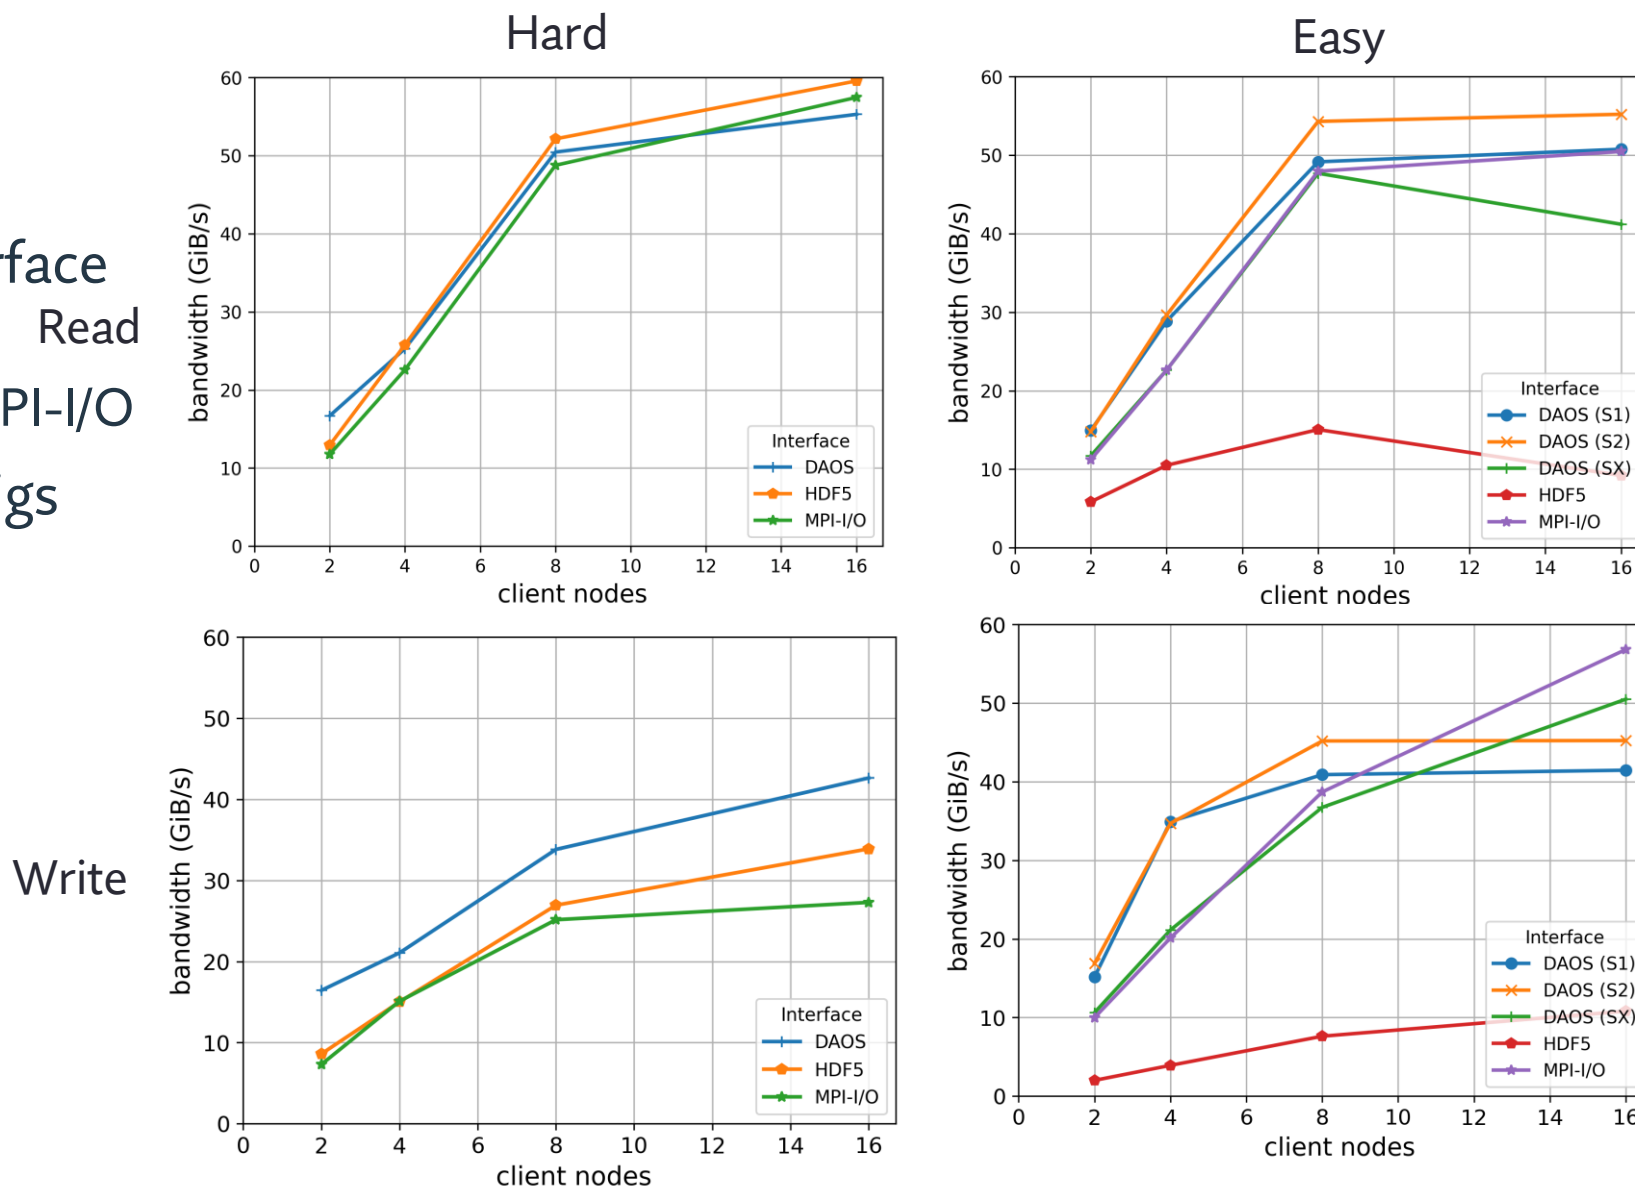
- DAOS provides “legacy” approaches
 - Requires some config to get high performance (i.e. number of workers/mount points)
 - <https://github.com/johannlombardi/daos-kernel>

```
mkdir /tmp/my_filesystem  
dfuse -m /tmp/my_filesystem --pool tutorial --cont adrians
```







...

```
fusermount3 -u /tmp/my_filesystem
```

- Benchmarking interface
 - DAOS api
 - Fuse for HDF5/MPI-I/O
- Wide range of configs



FUSE vs DAOS




DAOS Objects

- DAOS Object Types:
 - DAOS KV
 - Simplifies the interface to just give what is needed for key-value interactions
 - i.e. put/get
 - DAOS multi-level KV:
 - Extended key functionality
 - Two parts to allow distribution and data to be separately managed
 - distribution (dkey): All entries under the same dkey are guaranteed to be collocated on the same target
 - attribute (akey): Enables multiple values to be stored under the same key (or can be single value)
 - DAOS Array
 - Single dimension array referenced by a DAOS key
 - Specialised version of the multi-level KV, dkey and akey managed for you


DAOS objects

I/O Middleware View




Mapping

DAOS Layout View



Object

128-bit
object identifier



KVS

Composite key



Distribution key
Variable length
Integer, string, data structure, ...
Hashed/distributed across engines

Attribute key
Variable length
Integer, string, data structure, ...
Co-located on same engines
All akeys can be retrieved in a single RPC

Index
64-bit Integer
Address array of value
Support operation over range
of indices (= extent)

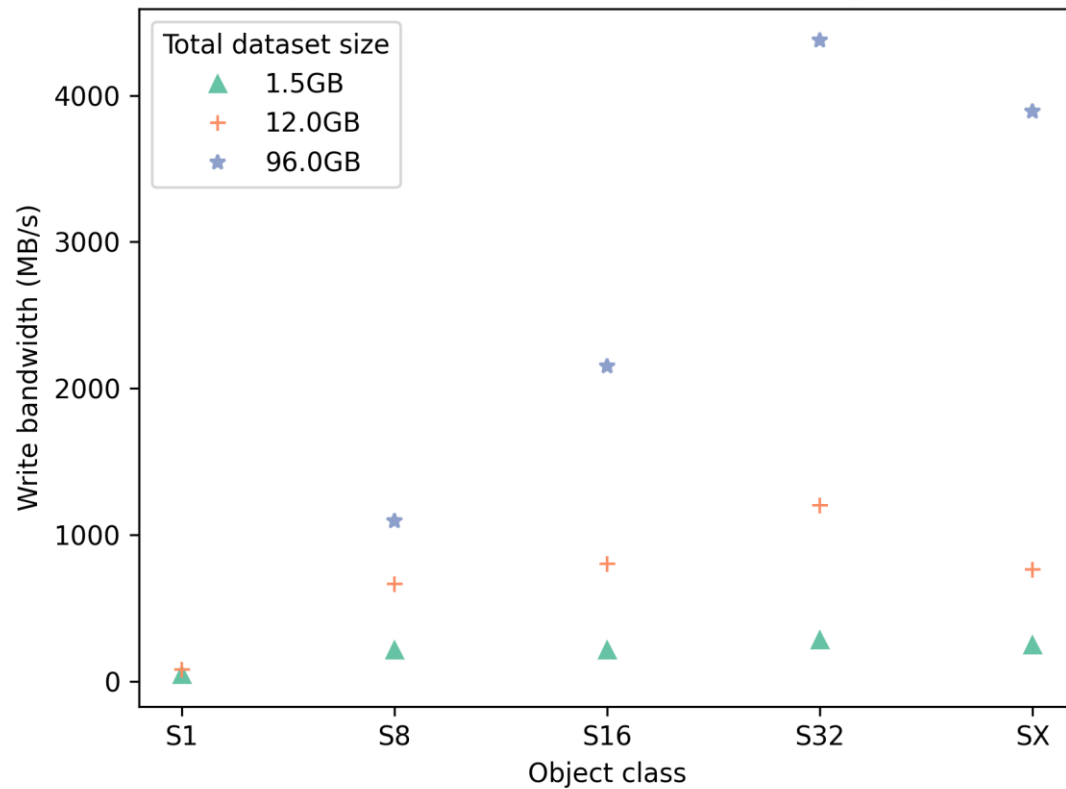
Record
Variable length value
From 1 byte to GB's

Object classes


- Object ID 128-bit space:
 - Lower 96 bits set by user
 - Unique OID allocator available in API for convenience
 - OID Embeds:
 - Object type
 - Object class (redundancy level and type – Replication, EC, None)
- DAOS allows configuration of *object class*
 - Sharding: Think Lustre striping
 - S1, S2, S4, S8, ..., S32, SX
 - Replication + Sharding: Think Lustre striping and object copying
 - RP_2G1, RP_2G2, RP_2G4, RP_2G8, ..., RP_2G32, RP_2GX
 - Can vary the replication higher than 2 (first number)
 - Erasure coding + Replication + Sharding:
 - OC_EC_2P1G1, OC_EC_2P1G2, ..., OC_EC_2P1GX
 - Can also vary the replication

Object class

benchio Performance on DAOS using 16 client nodes (48 processes per node) and 8 server nodes (2 engines per node)

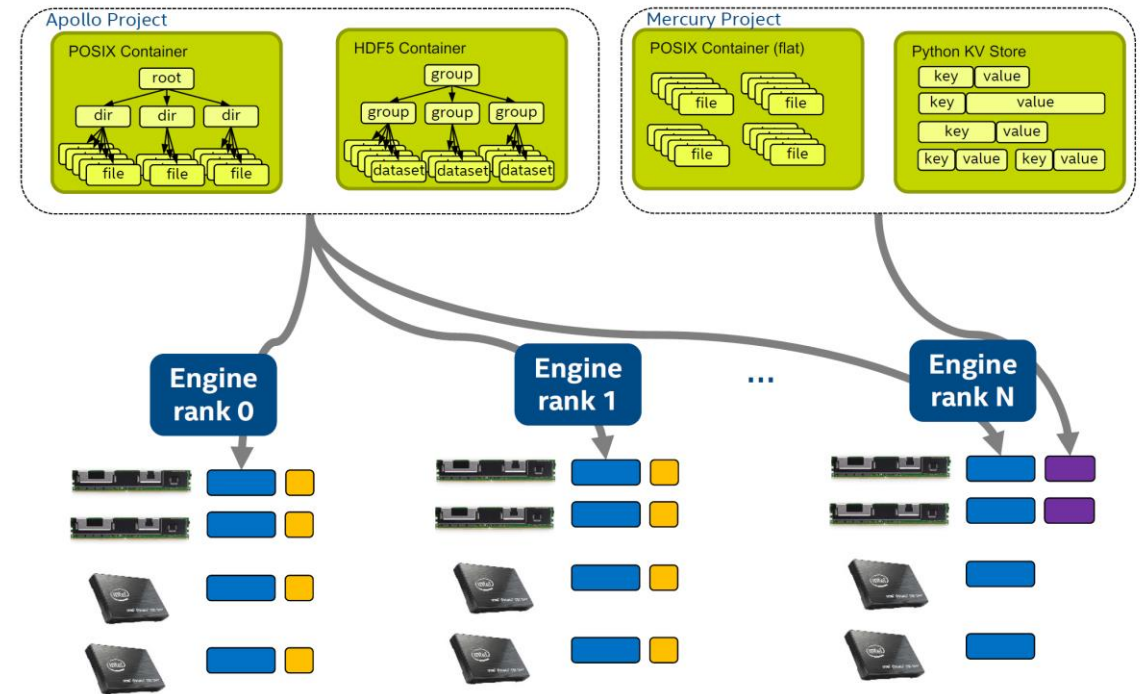


benchio Performance on DAOS using 4 client nodes (48 processes per node) and 8 server nodes (2 engines per node)



System Hierarchy

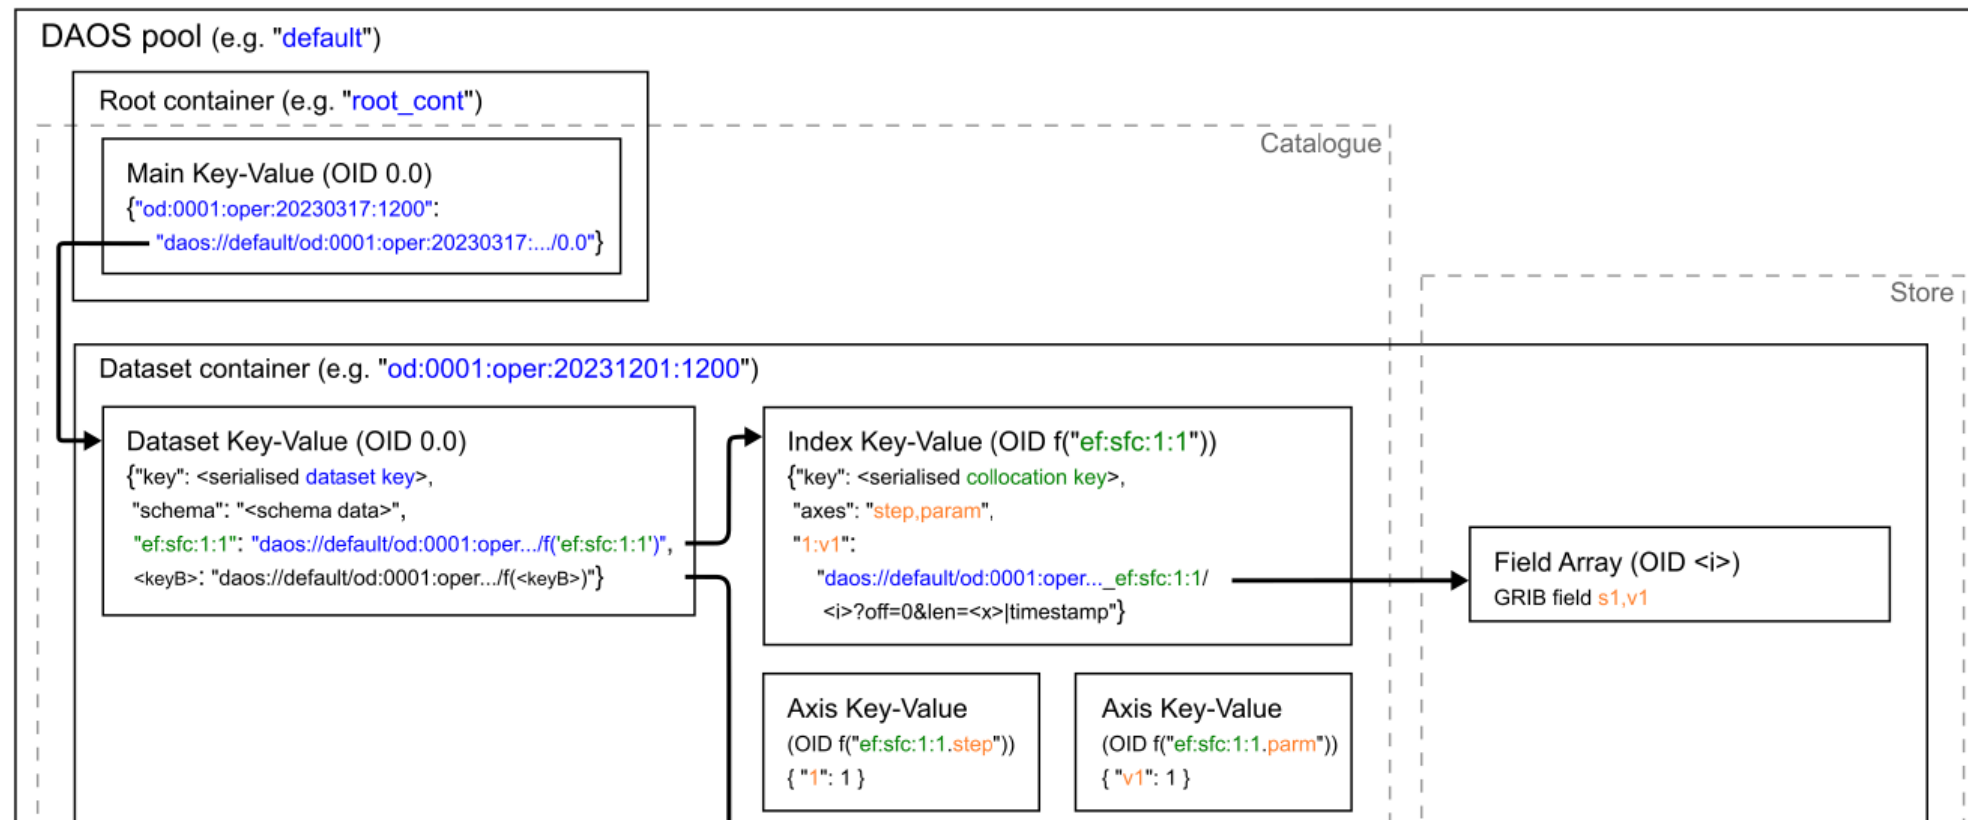
- Pools
 - Highest level construct in the system
 - Possibly analogous to a filesystem
 - Maybe more like a root of a tree
 - Maps to servers/storage hardware
 - Orders of ones per system
- Containers
 - Namespace/logical space inside a pool
 - Group together objects
 - Orders of hundred per pool
- Can specify configurations on both pools and containers
 - Access control, replication, redundancy, etc...
- Objects can inherit configurations from containers
 - And can define their own requirements
 - Any number



Example DAOS configuration/usage

```
FDB.archive(  
  Key({"class": "od", "expver": "0001", "stream": "oper", "date": "20231201", "time": "1200",  
      "type": "ef", "levtype": "sfc", "number": 1, "levelist": 1, "step": 1, "param": "v1"}),  
  buffer, length)
```

- Dataset key
- Collocation key
- Element key



DAOS API programming

```
rc = daos_init();
daos_pool_info_t pool_info;
daos_cont_info_t co_info;
rc = daos_pool_connect(o->pool, o->group, DAOS_PC_RW, &poh,
&pool_info, NULL);
rc = daos_cont_open(poh, o->cont, DAOS_COO_RW, &coh, &co_info, NULL);
rc = dfs_mount(poh, coh, O_RDWR, &dfs);
rc = dfs_write(dfs, obj, &sgl, off, NULL);
rc = dfs_read(dfs, obj, &sgl, off, &ret, NULL);
rc = dfs_umount(dfs);
rc = daos_cont_close(coh, NULL);
rc = daos_cont_destroy(poh, o->cont, 1, NULL);
rc = daos_pool_disconnect(poh, NULL);
rc = daos_fini();
```

DAOS C API - Pools

- Initialise and finalise the API

```
int daos_init(void);
```

```
int daos_fini(void);
```

- Connect to an existing pool

```
int daos_pool_connect(const char *pool, const char *sys, unsigned  
int flags, daos_handle_t *poh, daos_pool_info_t *info, daos_event_t  
*ev);
```

- Disconnect from a pool to ensure all I/O has concluded

```
int daos_pool_disconnect(daos_handle_t poh, daos_event_t *ev);
```

- Large scale parallel programs want to avoid every client opening the pool

- Costly operation at the server side
- Share the pool handle opened by a single client via MPI

```
daos_pool_local2global
```

```
daos_pool_global2local
```

Containers

- Can manually create/destroy containers at the command line

```
daos cont create mypool --label=mycont
```

```
Container UUID : 5d33d6e0-6c8b-4bf5-bb49-c8723bf30c91  
Container Label: mycont  
Container Type : unknown
```

```
Successfully created container 5d33d6e0-6c8b-4bf5-bb49-c8723bf30c91
```

- Can do this programmatically as well:

```
int daos_cont_create_with_label(daos_handle_t poh, const char *label, daos_prop_t *cont_prop,  
                                uuid_t *uuid, daos_event_t *ev);
```

```
int daos_cont_destroy(daos_handle_t poh, const char *cont, int force, daos_event_t *ev);
```

- Need to open/close the container from the program

```
int daos_cont_open(daos_handle_t poh, const char *cont, unsigned int flags, daos_handle_t *coh,  
                   daos_cont_info_t *info, daos_event_t *ev);
```

```
int daos_cont_close(daos_handle_t coh, daos_event_t *ev);
```

- As with the pool, want to avoid all clients opening manually

```
daos_cont_local2global
```

```
daos_cont_global2local
```

Minimal example

```
#include <daos.h>
int main(int argc, char **argv) {
    daos_handle_t poh, coh;
    daos_init();
    daos_pool_connect("mypool", NULL, DAOS_PC_RW, &poh, NULL, NULL);
    daos_cont_create_with_label(poh, "mycont", NULL, NULL, NULL);
    daos_cont_open(poh, "mycont", DAOS_COO_RW, &coh, NULL, NULL);
    /** do things */
    daos_cont_close(coh, NULL);
    daos_pool_disconnect(poh, NULL);
    daos_fini();
    return 0;
}
```


Arrays

- Current DAOS supports 1-D arrays
 - Future support for multi-dimensional arrays
- Create/open/close arrays

```
int daos_array_create(daos_handle_t coh, daos_obj_id_t oid, daos_handle_t th, daos_size_t cell_size, daos_size_t chunk_size, daos_handle_t *oh, daos_event_t *ev);
```

```
int daos_array_open(daos_handle_t coh, daos_obj_id_t oid, daos_handle_t th, unsigned int mode, daos_size_t *cell_size, daos_size_t *chunk_size, daos_handle_t *oh, daos_event_t *ev);
```

```
int daos_array_open_with_attr(daos_handle_t coh, daos_obj_id_t oid, daos_handle_t th, unsigned int mode, daos_size_t cell_size, daos_size_t chunk_size, daos_handle_t *oh, daos_event_t *ev);
```

```
int daos_array_close(daos_handle_t oh, daos_event_t *ev);
```

```
int daos_array_destroy(daos_handle_t oh, daos_handle_t th, daos_event_t *ev);
```

- `th` is a transaction handle
- `ev` is a completion event hand
- Both can be NULL

Arrays

- Chunk size: The striping chunk size
 - When to move on to another dkey
- Cell size: Size of the datatype of the array in bytes
- Open call will do create if it doesn't exist
 - This can be expensive to check
 - Open with attributes is cheaper if you know it exists
 - Bypasses checks (i.e. cell and chunk sizes)
- Data operations:

```
int daos_array_read(daos_handle_t oh, daos_handle_t th, daos_array_iod_t *iod,  
d_sg_list_t *sgl, daos_event_t *ev);
```

```
int daos_array_write(daos_handle_t oh, daos_handle_t th, daos_array_iod_t *iod,  
d_sg_list_t *sgl, daos_event_t *ev);
```

Arrays

- Specify range of data to read/write and location of data source/destination

```
daos_array_iod_t *iod, d_sg_list_t *sgl
```

- `iod` specifies the location of the data in the DAOS array
 - `sgl` specifies the location of the data in memory
- Multi-process writing/reading to DAOS array enabled
 - Can specify offsets and access patterns using the `iod`
- Support functions also available:

```
int daos_array_get_size(daos_handle_t oh, daos_handle_t th, daos_size_t *size, daos_event_t *ev);
```

```
int daos_array_set_size(daos_handle_t oh, daos_handle_t th, daos_size_t size, daos_event_t *ev);
```

```
int daos_array_get_attr(daos_handle_t oh, daos_size_t *chunk_size, daos_size_t *cell_size);
```

Arrays example

```
/* create array - if array exists just open it */
daos_array_create(coh, oid, DAOS_TX_NONE, 1, 1048576, &array, NULL);
daos_array_iod_t iod;
d_sg_list_t sgl;
daos_range_t rg;
d_iov_t iov;
/* set array location */
iod.arr_nr = 1; /** number of ranges / array iovec */
rg.rg_len = BUFLen; /** length */
rg.rg_idx = rank * BUFLen; /** offset */
iod.arr_rgs = &rg;
/* set memory location, each rank writing BUFLen */
sgl.sg_nr = 1;
d_iov_set(&iov, buf, BUFLen);
sgl.sg_iovs = &iov;
daos_array_write(array, DAOS_TX_NONE, &iod, &sgl, NULL);
daos_array_read(array, DAOS_TX_NONE, &iod, &sgl, NULL);
daos_array_close(array, NULL);
```

Key-Value API

```
int daos_kv_open(daos_handle_t coh, daos_obj_id_t oid, unsigned int mode, daos_handle_t *oh,
daos_event_t *ev);

int daos_kv_put(daos_handle_t oh, daos_handle_t th, uint64_t flags, const char *key, daos_size_t
size, const void *buf, daos_event_t *ev);

int daos_kv_get(daos_handle_t oh, daos_handle_t th, uint64_t flags, const char *key, daos_size_t
*size, void *buf, daos_event_t *ev);

int daos_kv_remove(daos_handle_t oh, daos_handle_t th, uint64_t flags, const char *key,
daos_event_t *ev);

int daos_kv_list(daos_handle_t oh, daos_handle_t th, uint32_t *nr, daos_key_desc_t *kds,
d_sg_list_t *sgl, daos_anchor_t *anchor, daos_event_t *ev);

int daos_kv_close(daos_handle_t oh, daos_event_t *ev);

int daos_kv_destroy(daos_handle_t oh, daos_handle_t th, daos_event_t *ev);
```

- Open returns a key handle, creates the key if required
 - Flag can be: DAOS_OO_RO, DAOS_OO_RW
- Put and Get work with the value associated with the key handle (oh)

Key-Value example

```
oid.hi = 0;
oid.lo = 1;
daos_obj_generate_oid(coh, &oid, DAOS_OF_KV_FLAT, 0, 0, 0);
daos_kv_open(coh, oid, DAOS_OO_RW, &kv, NULL);
/* set val buffer and size */
daos_kv_put(kv, DAOS_TX_NONE, 0, "key1", val_len1, val_buf1, NULL);
daos_kv_put(kv, DAOS_TX_NONE, 0, "key2", val_len2, val_buf2, NULL);
/* to fetch, can query the size first if not known */
daos_kv_get(kv, DAOS_TX_NONE, 0, "key1", &size, NULL, NULL);
get_buf = malloc (size);
daos_kv_get(kv, DAOS_TX_NONE, 0, "key1", &size, get_buf, NULL);
daos_kv_close(kv, NULL);
```

Multi-level KV

- Array and KV are specialised version
 - Multi-level provides full functionality

```
int daos_obj_open(daos_handle_t coh,  
daos_obj_id_t oid, unsigned int mode,  
daos_handle_t *oh, daos_event_t *ev);
```

```
int daos_obj_close(daos_handle_t oh,  
daos_event_t *ev);
```

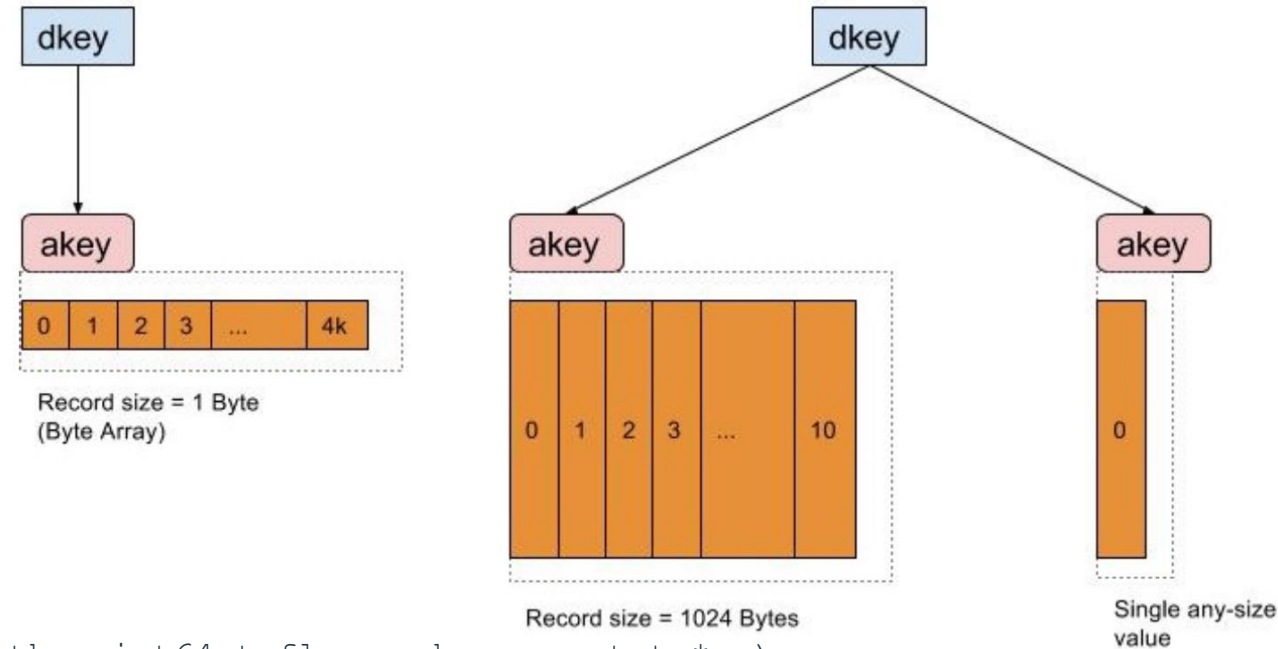
```
int daos_obj_punch(daos_handle_t oh, daos_handle_t th, uint64_t flags, daos_event_t *ev);
```

```
int daos_obj_punch_dkeys(daos_handle_t oh, daos_handle_t th, uint64_t flags, unsigned int nr, daos_key_t  
*dkeys, daos_event_t *ev);
```

```
int daos_obj_punch_akeys(daos_handle_t oh, daos_handle_t th, uint64_t flags, daos_key_t *dkey, unsigned int nr,  
daos_key_t *akeys, daos_event_t *ev);
```

```
int daos_obj_update(daos_handle_t oh, daos_handle_t th, uint64_t flags, daos_key_t *dkey, unsigned int nr,  
daos_iod_t *iods, d_sg_list_t *sgls, daos_event_t *ev);
```

```
int daos_obj_fetch(daos_handle_t oh, daos_handle_t th, uint64_t flags, daos_key_t *dkey, unsigned int  
nr, daos_iod_t *iods, d_sg_list_t *sgls, daos_iom_t *ioms, daos_event_t *ev);
```



Multi-level KV example

```
daos_obj_open(coh, oid, DAOS_OO_RW, &oh, NULL);
d_iov_set(&dkey, "dkey1", strlen("dkey1"));
d_iov_set(&sg_iov, buf, BUFLen);
sgl[0].sg_nr = 1;
sgl[0].sg_iovs = &sg_iov;
sgl[1].sg_nr = 1;
sgl[1].sg_iovs = &sg_iov;
d_iov_set(&iod[0].iod_name, "akey1", strlen("akey1"));
d_iov_set(&iod[1].iod_name, "akey2", strlen("akey2"));
iod[0].iod_nr = 1;
iod[0].iod_size = BUFLen;
iod[0].iod_recxs = NULL;
iod[0].iod_type = DAOS_IOD_SINGLE;
iod[1].iod_nr = 1;
iod[1].iod_size = 1;
recx.rx_nr = BUFLen;
recx.rx_idx = 0;
iod[1].iod_recxs = &recx;
iod[1].iod_type = DAOS_IOD_ARRAY;
daos_obj_update(oh, DAOS_TX_NONE, 0, &dkey, 2, &iod, &sgl, NULL);
```


Fortran interfacing

```
call daos_initialise(pool_name_c, cartcomm)
call daos_write_array(ndim, arraysize, arraygsize, arraysubsize, arraystart, out_data, object_class_c, blocksize, check_data, daosconfig, cartcomm)
call daos_write_object(ndim, arraysize, arraygsize, arraysubsize, arraystart, out_data, object_class_c, blocksize, check_data, daosconfig, cartcomm)
call daos_read_array(ndim, arraysize, arraygsize, arraysubsize, arraystart, read_data, object_class_c, daosconfig, cartcomm)
call daos_read_object(ndim, arraysize, arraygsize, arraysubsize, arraystart, read_data, object_class_c, daosconfig, cartcomm)
call daos_finish(iocomm)

array_obj_id.hi = 0;
array_obj_id.lo = 0;

uuid_generate_md5(array_uuid, seed, array_name, strlen(array_name));

memcpy(&(array_obj_id.hi), &(array_uuid[0]) + sizeof(uint64_t), sizeof(uint64_t));
memcpy(&(array_obj_id.lo), &(array_uuid[0]), sizeof(uint64_t));

daos_array_generate_oid(container_handle, &array_obj_id, DAOS_OT_ARRAY_BYTE, array_obj_class, 0, 0);

ierr = daos_array_open(container_handle, array_obj_id, DAOS_TX_NONE, DAOS_OO_RW, &cell_size, &local_block_size, &array_handle, NULL);

total_size = sizeof(double);
for(i=0; i<num_dims; i++){
    total_size = total_size * arraysusbsize[i];
}

iod.arr_nr = 1;
rg.rg_len = total_size;
rg.rg_idx = 0;
iod.arr_rgs = &rg;


sgl.sg_nr = 1;
d_iov_set(&iiov, &output_data[0], total_size);
sgl.sg_iovs = &iiov;

ierr = daos_array_read(array_handle, DAOS_TX_NONE, &iod, &sgl, NULL);


ierr = daos_array_destroy(array_handle, DAOS_TX_NONE, NULL);
```

Evaluate performance/approach

Access pattern A, writers,




Access pattern A, readers,




Optimised performance

Access pattern A, writers,




Access pattern A, readers,




Profiling

- Example breakdown of where time is being spent
 - Manual profiling

fdb-hammer/DAOS write bottlenecks
12 server nodes, 20 client nodes, 32 processes per client node
average per-process fdb-hammer wall-clock time: 27.5 seconds



fdb-hammer/DAOS read bottlenecks
12 server nodes, 20 client nodes, 32 processes per client node
average per-process fdb-hammer wall-clock time: 21.8 seconds



Approach/recommendations

- Key-Value contention
- For a specific benchmark run configured with contention across processes on indexing Key-Values:
 - 20 GiB/s write
 - 13 GiB/s read
- Tweaking the benchmark configuration to have all processes operate on a separate Key-Values:
 - 35 GiB/s write
 - 68 GiB/s read
- This may not be trivial or possible for all applications, but if design can achieve it then this improves performance

Approach/recommendations

- Avoid communications on/with the server where possible
- Cache objects locally in DRAM if possible
- Use `daos_array_open_with_attr` to avoid `daos_array_create` calls
 - Only supported for `DAOS_OT_ARRAY_BYTE`, not for `DAOS_OT_ARRAY`
 - Warning: the cell size and chunk size attributes need to be provided consistently on any future `daos_array_open_with_attr` to avoid data corruption
- `daos_array_get_size` calls can be expensive
 - Can store array size in our indexing Key-Values
 - Can manually calculate
 - Also possible to infer the size by reading with overallocation:
 - use `DAOS_OT_ARRAY_BYTE`, over-allocate the read buffer, and read without querying the size. The actual read size (`short_read`) will be returned
- `daos_cont_alloc_oids` is expensive, call it just once per writer process
 - Required to generate object ideas to use in calls but can generate many at one

Approach/recommendation

- Creating several containers (starting at ~300) in a DAOS pool reduces performance
- Opening the same container from all processes is expensive
 - this happens even if only a few containers exist in the DAOS pool
 - e.g. out of 20 seconds taken by a process to write 2000 fields, 1.5 seconds were spent just to open one container
 - we observed this starting at ~200 parallel processes
 - Sharing handles using MPI is the way to fix this
- Opening more than one container per process is very expensive
 - e.g. out of 30 seconds taken by a process to read 2000 fields, 6 seconds were spent just to open two containers

Approach/recommendations

- `daos_key_value_list` is expensive
- `daos_array_open_with_attrs`, `daos_kv_open` and `daos_array_generate_oid` are very cheap (no RPC)
- Normal `daos_array_open` is expensive
- `daos_cont_alloc_oids` is expensive
- `daos_kv_put` and `_get` are generally cheap
 - Value size impacts this
- `daos_obj_close`, `daos_cont_close` and `daos_pool_disconnect` are cheap
- Server configuration to use available networks/sockets/etc... important for performance
 - Just like any storage system or application

DAOS usage design

- Mapping data structures to KV and Array objects is key to getting good performance functionality
- We suggest mapping contiguous chunks of arrays to be stored to single DAOS array object
 - Collect multiple arrays with associated KV to make the whole array
- Can be as extreme as having a single value per KV
 - Significant overheads in this
- Depends on your application data structures you may want to aggregate less data for I/O
 - Group based on meaningful/scientific dimensions
- HDF5 or similar hierarchies could map well to Keys with Arrays

Summary

- DAOS provides high performance storage
 - 90+ GB/s per server is possible
 - Hardware and configuration dependent, just like all I/O
- Built in replication and redundancy under you/user control
- Different interfaces available
 - Filesystem for zero cost porting
 - Simple file like access for slightly improved performance at little effort
 - DAOS APIs for full functionality
- DAOS interface enables changing I/O granularity/patterns for bigger benefits

