

Tutorial on the DAOS API

SC24 Tutorial

Mohamad Chaarawi



Notices and Disclaimers

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration.

No product or component can be absolutely secure.

Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. For more complete information about performance and benchmark results, visit <http://www.intel.com/benchmarks>.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. For more complete information visit <http://www.intel.com/benchmarks>.

Intel Advanced Vector Extensions (Intel AVX) provides higher throughput to certain processor operations. Due to varying processor power characteristics, utilizing AVX instructions may cause a) some parts to operate at less than the rated frequency and b) some parts with Intel® Turbo Boost Technology 2.0 to not achieve any or maximum turbo frequencies. Performance varies depending on hardware, software, and system configuration and you can learn more at <http://www.intel.com/go/turbo>.

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Cost reduction scenarios described are intended as examples of how a given Intel-based product, in the specified circumstances and configurations, may affect future costs and provide cost savings. Circumstances will vary. Intel does not guarantee any costs or cost reduction.

Intel does not control or audit third-party benchmark data or the web sites referenced in this document. You should visit the referenced web site and confirm whether referenced data are accurate.

© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.

DAOS User API

- Pools:

- Connect, disconnect, query

- Containers:

- Create, destroy, open, close, query, properties, attributes

- Objects:

- Access APIs based on type

DAOS Usage flow

- Initialize DAOS stack and connect to a Pool
- Create / open a container
- Access an object in the container through the unique oid
 - Open object, update/fetch/list, close object
 - Transaction API available
- Close container & disconnect from pool

Connect to Pool

- First thing you usually do in your program is initialize DAOS and connect to your pool:
 - `int daos_init(void);`
 - `int daos_pool_connect(const char *pool, const char *sys, unsigned int flags, daos_handle_t *poh, daos_pool_info_t *info, daos_event_t *ev);`
- When finished, disconnect from your pool and finalize daos:
 - `int daos_pool_disconnect(daos_handle_t poh, daos_event_t *ev);`
 - `int daos_fini(void);`
- In an MPI program, consider connecting only from 1 client, and sharing the pool handle (poh) using the pool l2g, g2l functions:
 - `daos_pool_local2global`
 - `daos_pool_global2local`

Creating a Container

■ Using the daos tool:

```
daos cont create mypool mycont
```

```
Container UUID : 5d33d6e0-6c8b-4bf5-bb49-c8723bf30c91  
Container Label: mycont  
Container Type : unknown
```

```
Successfully created container 5d33d6e0-6c8b-4bf5-bb49-c8723bf30c91
```

■ Using the API:

- `int daos_cont_create_with_label (daos_handle_t poh, const char *label, daos_prop_t *cont_prop, uuid_t *uuid, daos_event_t *ev);`
- `int daos_cont_destroy (daos_handle_t poh, const char *cont, int force, daos_event_t *ev);`

Accessing a Container

- Need to open a container to access object in it:
 - `int daos_cont_open(daos_handle_t poh, const char *cont, unsigned int flags, daos_handle_t *coh, daos_cont_info_t *info, daos_event_t *ev);`
- Close container when done:
 - `int daos_cont_close(daos_handle_t coh, daos_event_t *ev);`
- In an MPI program, consider opening only from 1 client, and sharing the container handle (coh) using the cont l2g, g2l functions:
 - `daos_cont_local2global`
 - `daos_cont_global2local`

Recap Program Flow

```
#include <daos.h>

int main(int argc, char **argv)
{
    daos_handle_t    poh, coh;

    daos_init();
    daos_pool_connect("mypool", NULL, DAOS_PC_RW, &poh, NULL, NULL);
    daos_cont_create_with_label(poh, "mycont", NULL, NULL, NULL);
    daos_cont_open(poh, "mycont", DAOS_COO_RW, &coh, NULL, NULL);

    /** do things */

    daos_cont_close(coh, NULL);
    daos_pool_disconnect(poh, NULL);
    daos_fini();
    return 0;
}
```


DAOS Object Types

■ DAOS Object Types:

- DAOS KV (put, get, list, remove)
 - 1 string key, 1 opaque value
- DAOS ARRAY
 - 1D array of fixed sized value
- DAOS Multi-Level KV (lower level)

■ Object ID 128-bit space:

- Lower 96 bits set by user
 - Unique OID allocator available in API for convenience
- OID Embeds:
 - Object type
 - Object class (redundancy level and type – Replication, EC, None)

DAOS KV Object

- KV store API that provides:

- Put, Get, Remove, List

- API:

- `int daos_kv_open(daos_handle_t coh, daos_obj_id_t oid, unsigned int mode, daos_handle_t *oh, daos_event_t *ev);`
- `int daos_kv_put(daos_handle_t oh, daos_handle_t th, uint64_t flags, const char *key, daos_size_t size, const void *buf, daos_event_t *ev);`
- `int daos_kv_get(daos_handle_t oh, daos_handle_t th, uint64_t flags, const char *key, daos_size_t *size, void *buf, daos_event_t *ev);`
- `int daos_kv_remove(daos_handle_t oh, daos_handle_t th, uint64_t flags, const char *key, daos_event_t *ev);`
- `int daos_kv_list(daos_handle_t oh, daos_handle_t th, uint32_t *nr, daos_key_desc_t *kds, d_sg_list_t *sgl, daos_anchor_t *anchor, daos_event_t *ev);`
- `int daos_kv_close(daos_handle_t oh, daos_event_t *ev);`
- `int daos_kv_destroy(daos_handle_t oh, daos_handle_t th, daos_event_t *ev);`

KV Conditional Operations

- By default, KV put/get operations do not check “existence” of key before operations:
 - Put on a key overwrites the value of that key
 - Get of a key does not fail if key does not exist, just returns 0 size.
 - Remove of a key does not fail if key does not exist.
- One can use conditional flags to achieve a different behavior:
- `DAOS_COND_KEY_INSERT`: Insert a key if it doesn't exist, fail if it does.
- `DAOS_COND_KEY_UPDATE`: Update a key if it exists, fail if it doesn't.
- `DAOS_COND_KEY_GET`: Get key value if it exists, fail if it doesn't.
- `DAOS_COND_KEY_REMOVE`: Remove a key if it exists, fail if it doesn't.

KV put/get example

```
/** init, connect, cont_open */

oid.hi = 0;
oid.lo = 1;
daos_obj_generate_oid(coh, &oid, DAOS_OF_KV_FLAT, 0, 0, 0);
daos_kv_open(coh, oid, DAOS_OO_RW, &kv, NULL);

/** set val buffer and size */
daos_kv_put(kv, DAOS_TX_NONE, 0, "key1", val_len1, val_buf1, NULL);
daos_kv_put(kv, DAOS_TX_NONE, 0, "key2", val_len2, val_buf2, NULL);

/** to fetch, can query the size first if not known */
daos_kv_get(kv, DAOS_TX_NONE, 0, "key1", &size, NULL, NULL);
get_buf = malloc (size);
daos_kv_get(kv, DAOS_TX_NONE, 0, "key1", &size, get_buf, NULL);
daos_kv_close(kv, NULL);

/** free buffer, cont_close, disconnect, finalize */
```

KV list example

```
/** enumerate keys in the KV */
daos_anchor_t    anchor = {0};
d_sg_list_t      sgl;
d_iov_t          sg_iov;

/** size of buffer to hold as many keys in memory */
buf = malloc(ENUM_DESC_BUF);
d_iov_set(&sg_iov, buf, ENUM_DESC_BUF);
sgl.sg_nr        = 1;
sgl.sg_nr_out    = 0;
sgl.sg_iovs      = &sg_iov;
```

```
daos_key_desc_t kds[ENUM_DESC_NR];

while (!daos_anchor_is_eof(&anchor)) {
    /** how many keys to attempt to fetch in one call */
    uint32_t      nr = ENUM_DESC_NR;

    memset(buf, 0, ENUM_DESC_BUF);
    daos_kv_list(kv, DAOS_TX_NONE, &nr, kds, &sgl,
                  &anchor, NULL);

    if (nr == 0)
        continue;

    /** buf now container nr keys */
    /** kds arrays has length of each key */
}
```

DAOS Array Object

- Array object to manage records:

- 1 Dimensional

- Array Management API:

- `int daos_array_create(daos_handle_t coh, daos_obj_id_t oid, daos_handle_t th, daos_size_t cell_size, daos_size_t chunk_size, daos_handle_t *oh, daos_event_t *ev);`
- `int daos_array_open(daos_handle_t coh, daos_obj_id_t oid, daos_handle_t th, unsigned int mode, daos_size_t *cell_size, daos_size_t *chunk_size, daos_handle_t *oh, daos_event_t *ev);`
- `int daos_array_close(daos_handle_t oh, daos_event_t *ev);`
- `int daos_array_destroy(daos_handle_t oh, daos_handle_t th, daos_event_t *ev);`

DAOS Array Access API

- Reading & writing record to an Array:
- `int daos_array_read(daos_handle_t oh, daos_handle_t th, daos_array_iod_t *iod, d_sg_list_t *sgl, daos_event_t *ev);`
- `int daos_array_write(daos_handle_t oh, daos_handle_t th, daos_array_iod_t *iod, d_sg_list_t *sgl, daos_event_t *ev);`
- `int daos_array_get_size(daos_handle_t oh, daos_handle_t th, daos_size_t *size, daos_event_t *ev);`
- `int daos_array_set_size(daos_handle_t oh, daos_handle_t th, daos_size_t size, daos_event_t *ev);`
- `int daos_array_get_attr(daos_handle_t oh, daos_size_t *chunk_size, daos_size_t *cell_size);`

DAOS Array example

```
/** create array - if array exists just open it */
daos_array_create(coh, oid, DAOS_TX_NONE, 1, 1048576, &array, NULL);

daos_array_iod_t iod;
d_sg_list_t      sgl;
daos_range_t     rg;
d_iov_t          iov;

/** set array location */
iod.arr_nr = 1; /** number of ranges / array iovec */
rg.rg_len = BUFLen; /** length */
rg.rg_idx = rank * BUFLen; /** offset */
iod.arr_rgs = &rg;

/** set memory location, each rank writing BUFLen */
sgl.sg_nr = 1;
d_iov_set(&iov, buf, BUFLen);
sgl.sg_iovs = &iov;

daos_array_write(array, DAOS_TX_NONE, &iod, &sgl, NULL);
daos_array_read(array, DAOS_TX_NONE, &iod, &sgl, NULL);
daos_array_close(array, NULL);
```




DAOS Multi-Level KV Object

- 2 level keys:

- Distribution Key - Dkey (collocate all entries under it), holds multiple akeys
- Attribute Key - Akey (lower level to address records)
- Both are opaque (support any size / type)

- Value types (under akey):

- Single value: one blob (traditional value in KV store)
- Array value:
 - 1 record size per akey
 - Array of records that can be updates via different extents / iovec



▪ **Intentionally very flexible, rich API; but at the expense of higher complexity for the regular user.**

Object Management API

- `int daos_obj_open(daos_handle_t coh, daos_obj_id_t oid, unsigned int mode, daos_handle_t *oh, daos_event_t *ev);`
- `int daos_obj_close(daos_handle_t oh, daos_event_t *ev);`
- `int daos_obj_punch(daos_handle_t oh, daos_handle_t th, uint64_t flags, daos_event_t *ev);`
- `int daos_obj_punch_dkeys(daos_handle_t oh, daos_handle_t th, uint64_t flags, unsigned int nr, daos_key_t *dkeys, daos_event_t *ev);`
- `int daos_obj_punch_akeys(daos_handle_t oh, daos_handle_t th, uint64_t flags, daos_key_t *dkey, unsigned int nr, daos_key_t *akeys, daos_event_t *ev);`

Object IO API

```
■ int daos_obj_update(daos_handle_t oh, daos_handle_t th,  
    uint64_t flags, daos_key_t *dkey, unsigned int nr,  
    daos_iod_t *iods, d_sg_list_t *sgls, daos_event_t *ev);
```



```
daos_key_t iod_name; /* akey */  
daos_iod_type_t iod_type; /* value type (single value or array value) */  
daos_size_t iod_size; /* value or record size */  
uint32_t iod_nr; /* number of extents (1 for SV) */  
daos_recx_t *iod_recxs; /* array of extents - offset, length pairs */  
  
uint32_t sg_nr;  
uint32_t sg_nr_out;  
d_iov_t *sg_iovs;
```

```
■ int daos_obj_fetch(daos_handle_t oh, daos_handle_t th,  
    uint64_t flags, daos_key_t *dkey, unsigned int nr,  
    daos_iod_t *iods, d_sg_list_t *sgls, daos_iom_t *ioms,  
    daos_event_t *ev);
```

Object Enumerate API

```
int daos_obj_list_dkey(daos_handle_t oh, daos_handle_t  
th, uint32_t *nr, daos_key_desc_t *kds, d_sg_list_t  
*sgl, daos_anchor_t *anchor, daos_event_t *ev);
```

```
int daos_obj_list_akey(daos_handle_t oh, daos_handle_t  
th, daos_key_t *dkey, uint32_t *nr, daos_key_desc_t  
*kds, d_sg_list_t *sgl, daos_anchor_t *anchor,  
daos_event_t *ev);
```

DAOS Object Update Example

```
daos_obj_open(coh, oid, DAOS_OO_RW, &oh, NULL);
d_iov_set(&dkey, "dkey1", strlen("dkey1"));

d_iov_set(&sg_iov, buf, BUFLLEN);
sgl[0].sg_nr = 1;
sgl[0].sg_iovs = &sg_iov;
sgl[1].sg_nr = 1;
sgl[1].sg_iovs = &sg_iov;

d_iov_set(&iod[0].iod_name, "akey1", strlen("akey1"));
d_iov_set(&iod[1].iod_name, "akey2", strlen("akey2"));

iod[0].iod_nr = 1;
iod[0].iod_size = BUFLLEN;
iod[0].iod_recxs = NULL;
iod[0].iod_type = DAOS_IOD_SINGLE;

iod[1].iod_nr = 1;
iod[1].iod_size = 1;
recx.rx_nr = BUFLLEN;
recx.rx_idx = 0;
iod[1].iod_recxs = &recx;
iod[1].iod_type = DAOS_IOD_ARRAY;

daos_obj_update(oh, DAOS_TX_NONE, 0, &dkey, 2, &iod, &sgl, NULL);
```

DAOS Object Fetch Example

```
daos_obj_open(coh, oid, DAOS_OO_RW, &oh, NULL);
d_iov_set(&dkey, "dkey1", strlen("dkey1"));

d_iov_set(&sg_iov, buf, BUFLen);
sgl[0].sg_nr = 1;
sgl[0].sg_iovs = &sg_iov;
sgl[1].sg_nr = 1;
sgl[1].sg_iovs = &sg_iov;

d_iov_set(&iod[0].iod_name, "akey1", strlen("akey1"));
d_iov_set(&iod[1].iod_name, "akey2", strlen("akey2"));

iod[0].iod_nr = 1;
iod[0].iod_size = BUFLen; /** if size is not known, use DAOS_REC_ANY and NULL sgl */
iod[0].iod_recxs = NULL;
iod[0].iod_type = DAOS_IOD_SINGLE;

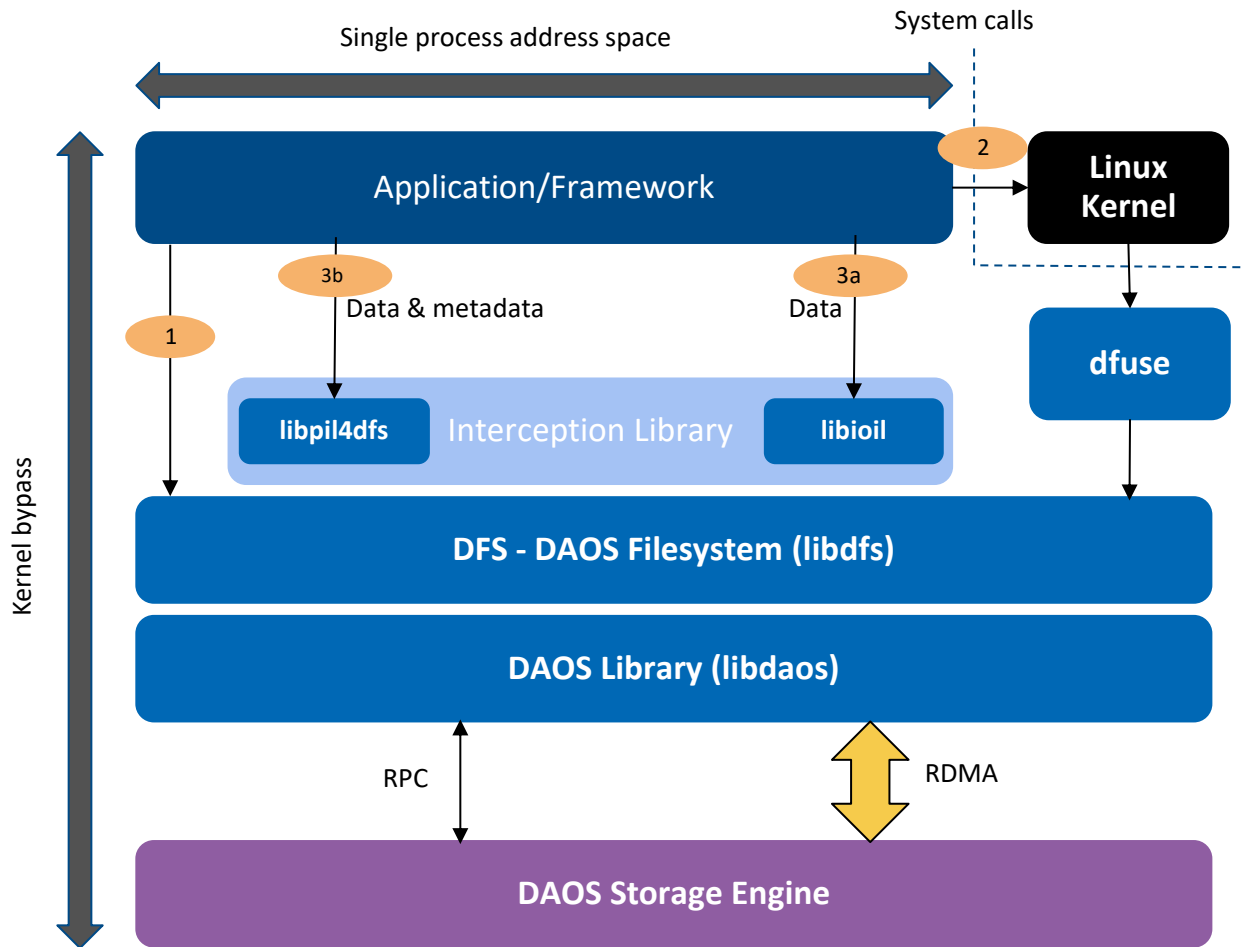
iod[1].iod_nr = 1;
iod[1].iod_size = 1; /** if size is not known, use DAOS_REC_ANY and NULL sgl */
recx.rx_nr = BUFLen;
recx.rx_idx = 0;
iod[1].iod_recxs = &recx;
iod[1].iod_type = DAOS_IOD_ARRAY;

daos_obj_fetch(oh, DAOS_TX_NONE, 0, &dkey, 2, &iod, &sgl, NULL, NULL);
```

More examples

- https://github.com/daos-stack/daos/blob/master/src/tests/simple_obj.c

POSIX Support & Interception



- 1 Userspace DFS library with API like POSIX
 - **Require** application changes
 - Low latency & high concurrency
 - No caching
 - 2 DFUSE daemon to support POSIX API
 - **No** application changes
 - VFS mount point & high latency
 - Caching by Linux kernel
 - 3 DFUSE + Interception library
 - **No** application changes
 - 2 flavors using LD_PRELOAD
- 3a libioil
 - (f)read/write interception
 - Metadata via dfuse
 - 3b libpil4dfs
 - Data & metadata interception
 - Aim at delivering same performance as #1 w/o any application change
 - Mmap & binary execution via fuse

How to use DFS?

- You should have access to a pool (identified by a string label).
- Create a POSIX container with the daos tool:
 - `daos cont create mypool mycont --type=POSIX`
 - Alternatively, you can programmatically create a container to use directly in your application (if you are using DFS and changing your app).
- Open the DFS mount:
 - `dfs_connect (mypool, mycont, O_RDWR, .. &dfs);`
 - `dfs_disconnect (dfs);`

DFS API

POSIX	DFS
mkdir(), rmdir()	dfs_mkdir(), dfs_rmdir()
open(), close(), access()	dfs_open(), dfs_release(),dfs_lookup()
pwritev(), preadv()	dfs_read/write()
{set,get,list,remove}xattr()	dfs_{set,get,list,remove}xattr
stat(), fstat()	dfs_stat(),ostat()
readdir() ...	dfs_readdir() ...

- Mostly 1-1 mapping from POSIX API to DFS API.
- Instead of File & Directory descriptors, use DFS objects.
- All calls need the DFS mount which is usually done once initialization.

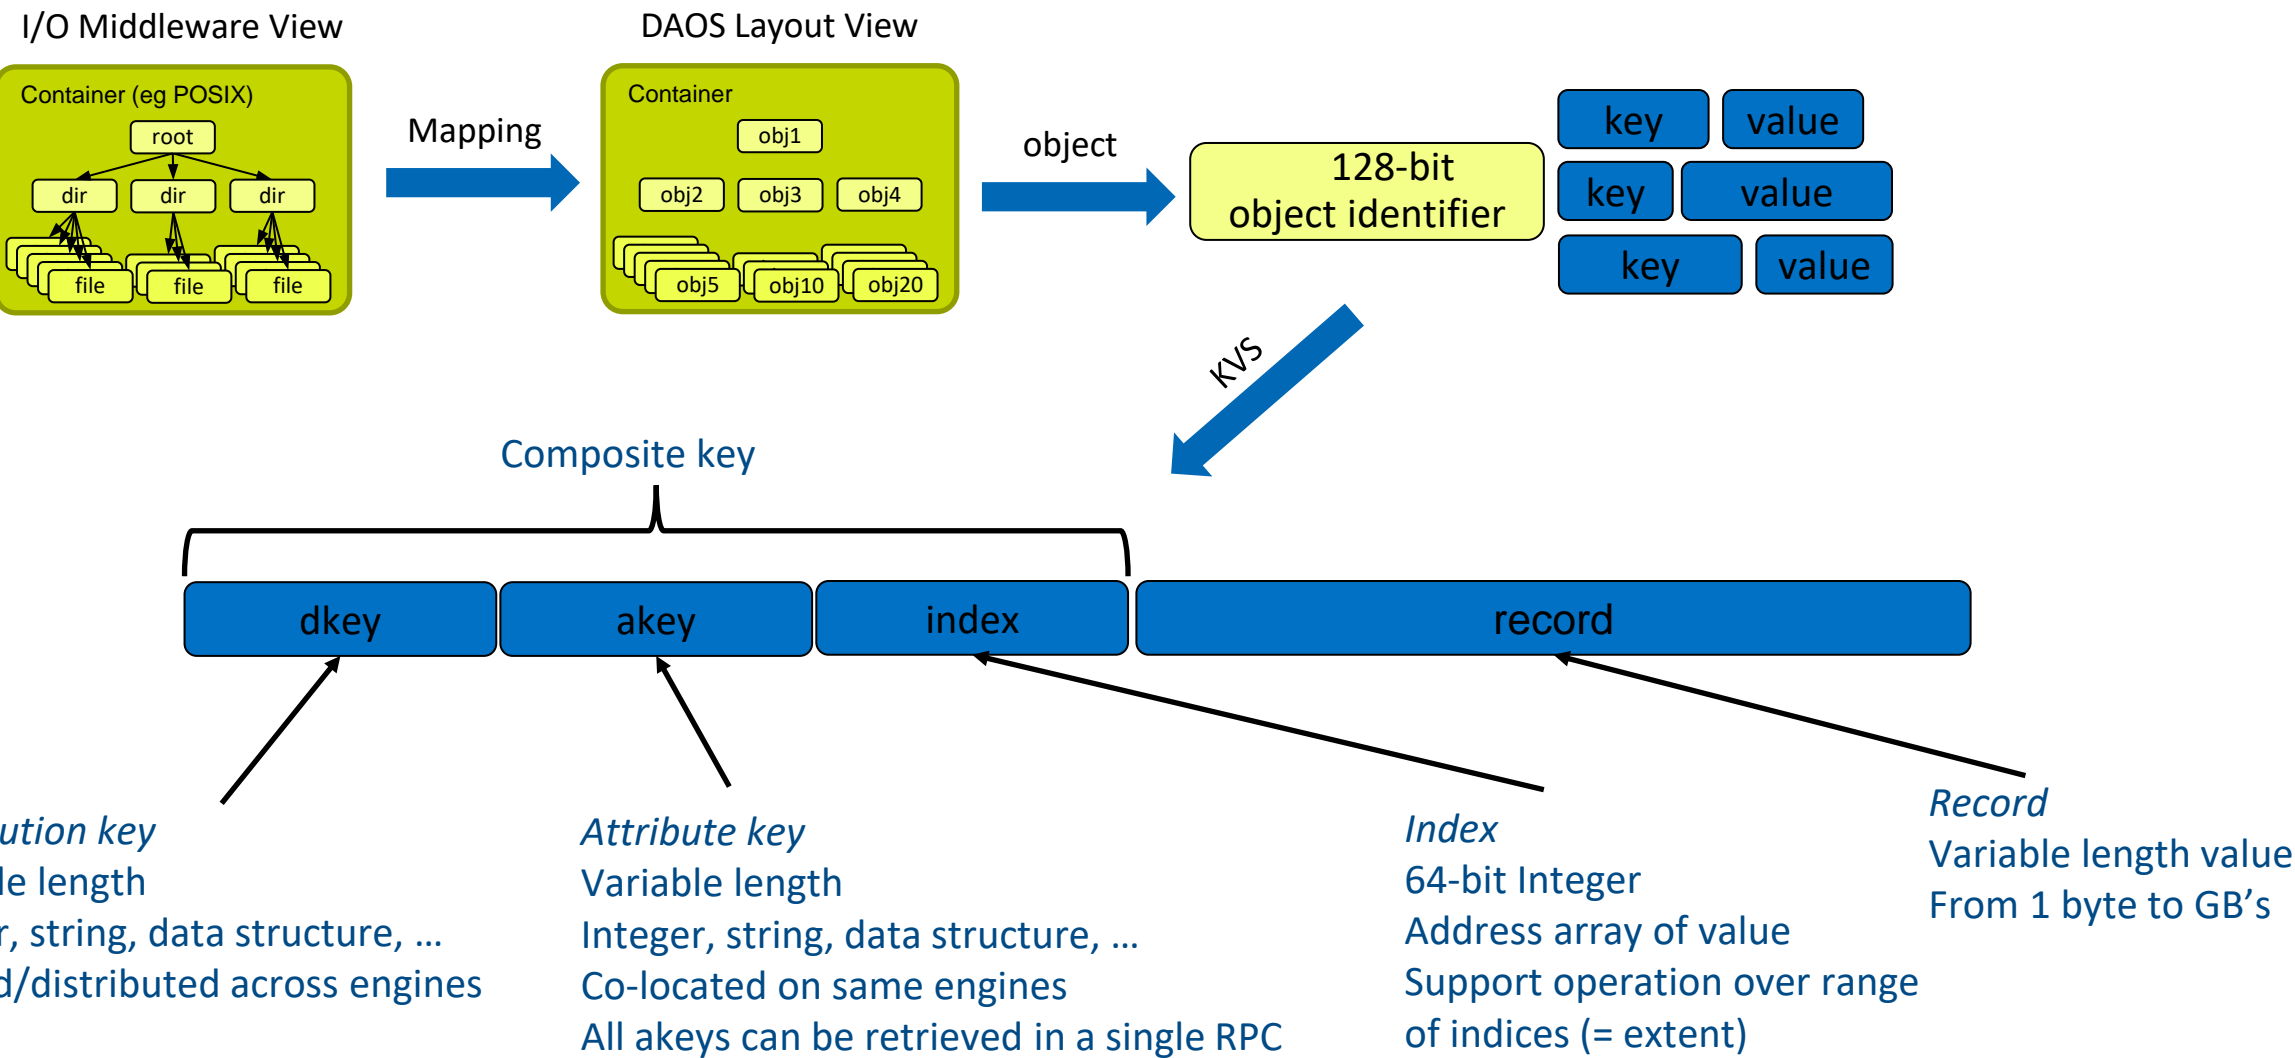
DFUSE

- To mount an existing POSIX container with dfuse, run the following command:
 - `dfuse mypool mycont -m /mnt/dfuse`
 - No one can access your container / mountpoint unless access is provided on the pool and container (through ACLs).
- Now you have a parallel file system under /mnt/dfuse on all nodes where that is mounted
 - Access files / directories as any namespace in the container, and applications can run without any modifications (the easy path).
- Interception Libraries:
 - This library works in conjunction with dfuse and allow to interception of POSIX I/O calls and issue the I/O operations directly from the application context through libdaos without any application changes.
 - This provides kernel-bypass for I/O. To use this set the LD_PRELOAD to point to the shared library in the DOAS install dir
 - `LD_PRELOAD=/path/to/daos/install/lib/libioil.so` or `libpil4dfs.so`

Hands On

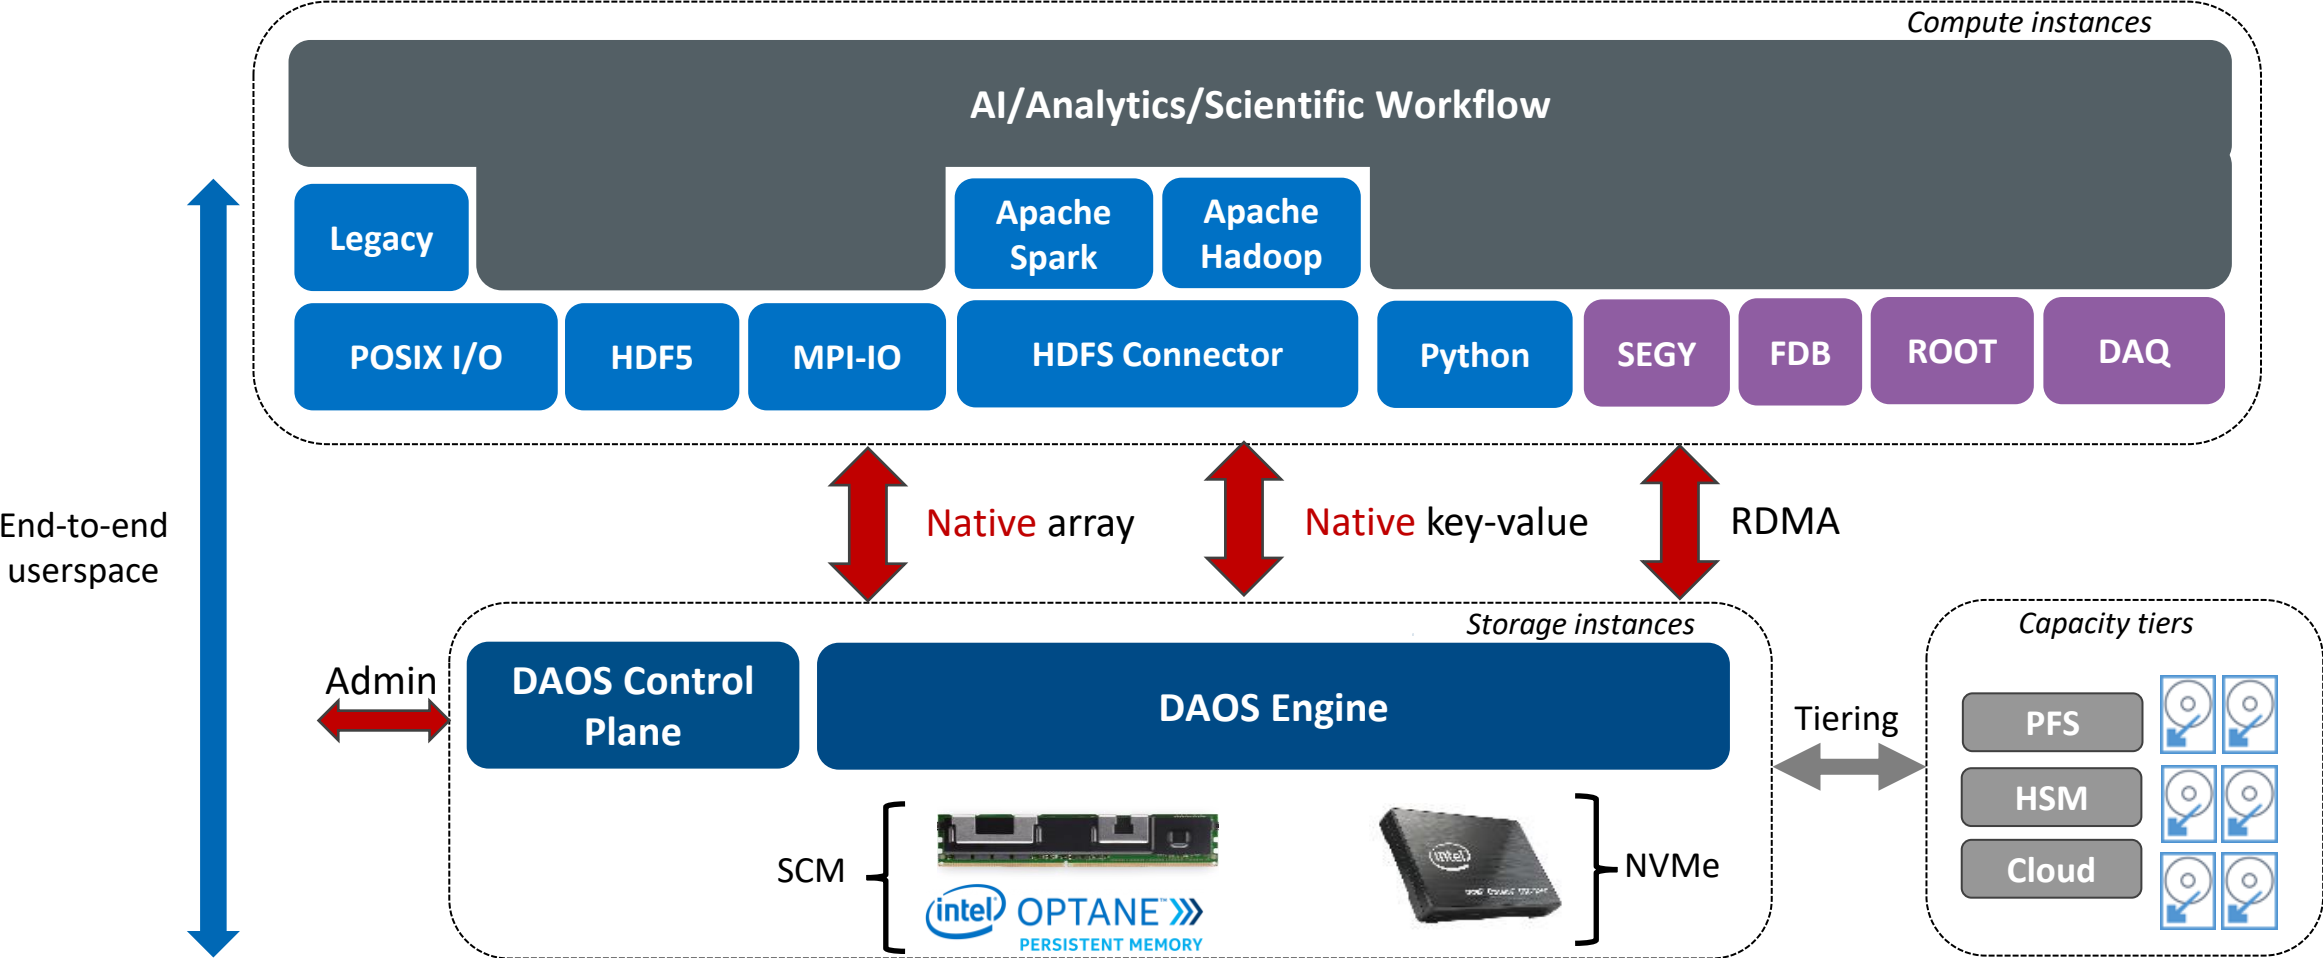


DAOS Data Model: Objects






DAOS Ecosystem


- Generic I/O middleware supported today
- Domain-specific data models under development in co-design with partners



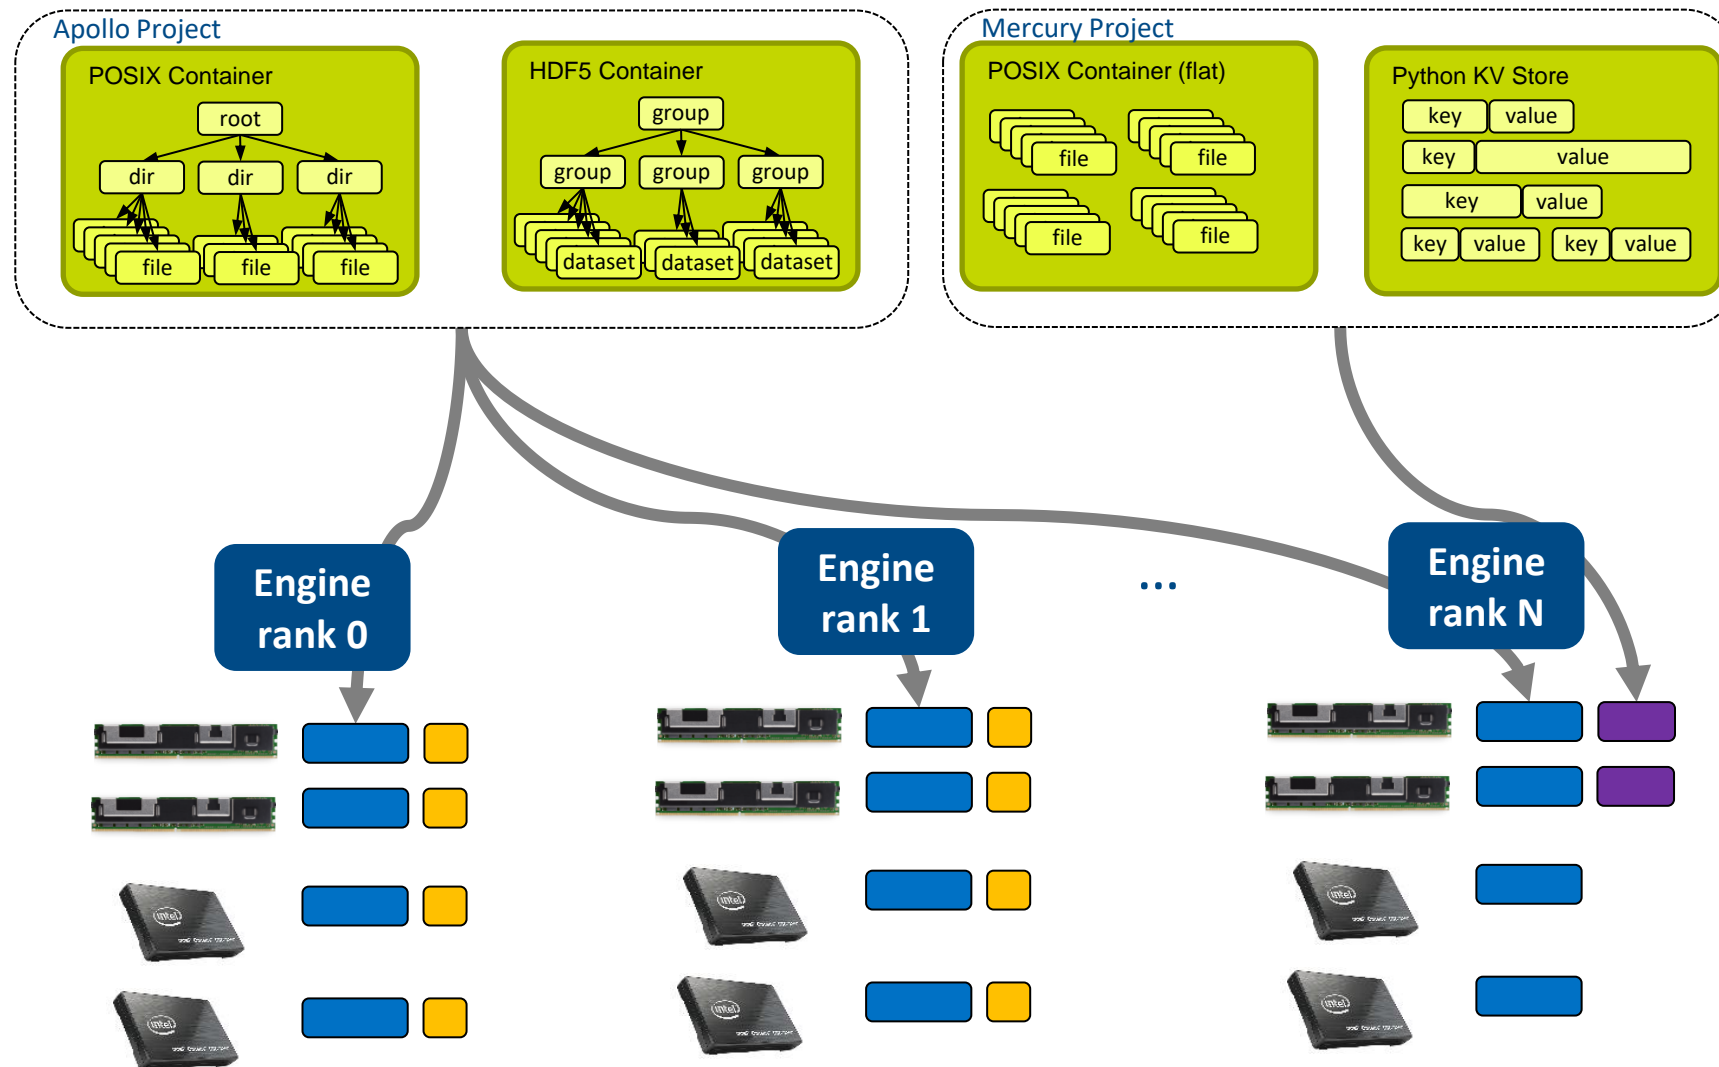
DAOS Data Model: Pools

Example:

Pool 1		Project Apollo	100PB usable	20TB/s	200M IOPS
Pool 2		Project Gemini	10PB usable	2TB/s	20M IOPS
Pool 3		Project Mercury	30TB usable	80GB/s	2M IOPS



DAOS Data Model: Container



Transactions

- Open/close transaction
 - Open returns transaction handle to use in object operations
 - Close just frees the handle (does not commit)
- Commit/abort/restart transaction
- Conflict detection:
 - WR, RW, WW conflicts

```
    daos_tx_open(coh, &th, ...);  
restart:  
    daos_obj_fetch(..., th, ...);  
    daos_obj_update(..., th, ...);  
    daos_obj_fetch(..., th, ...);  
    daos_obj_update(..., th, ...);  
    daos_obj_dkey_punch(..., th, ...);  
    rc = daos_tx_commit(th, ...);  
    if (rc == -DER_RESTART) {  
        daos_tx_restart(th, ...);  
        goto restart;  
    }  
    daos_tx_close(th, ...);
```